

# Documentation for the **sf1r-module**

Paolo D'Apice

March 30, 2012

## Abstract

This document contains the documentation for the Nginx **sf1r-module** enabling the communication with the **SF1**.

## 1 Overview

The objective of the **sf1r-module** is to allow the **SF1** to handle requests via HTTP, as shown in Figure 1 on the following page. The **sf1r-module** relies on the **libsflr** [1], which implements two driver clients for the **SF1**:

**single** allowing the connection to an **SF1** instance, hence implementing a direct HTTP front-end to the **SF1**

**distributed** allowing the connection to a cluster of **SF1** instances managed with ZooKeeper, acting as an active reverse proxy to the cluster

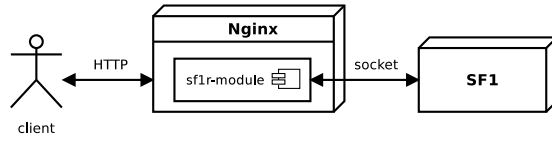
The diagram in Figure 2 on the next page shows how a single request is processed:

1. HTTP request received
2. **sf1r-module** request handler
3. request sent to **SF1**
4. response received from the **SF1**
5. response handler
6. HTTP response sent

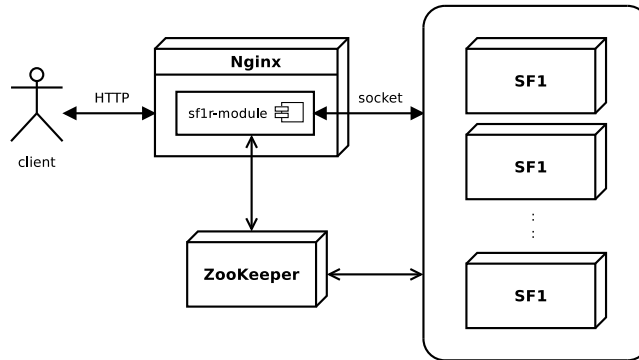
## Implementation Notes

**Logging** The **libsflr** uses the Google Logging Library. Within Nginx it is possible to partially control the logging behavior by properly defining environment variables, such as **GLDG\_log\_dir** or **GLDG\_minloglevel** [5]. However, there are still some open issues about the initialization/finalization of the logging system:

- logging from a library  
<http://code.google.com/p/google-glog/issues/detail?id=113>



(a) Single driver.



(b) Distributed driver.

Figure 1: Deployment diagram.




Figure 2: Communication diagram.

Listing 1: Sample HTTP request

```
HTTP POST /path/documents/search
content-type: application/json
{ "collection":"example", "header":{"check_time":true},
  "search":{"keywords":"america"}, "limit":10 }
```

- calling `ShutdownGoogleLogging()` after `InitGoogleLogging()` does not work  
<http://code.google.com/p/google-glog/issues/detail?id=83>

Such issues can be workarounded until they get fixed upstream simply by not explicitly initializing/finalizing the logging system.

**Threads** The ZooKeeper client used within `libsflr` spawns two thread: one for I/O and one for event handling. Nginx implementation, on the other hand, is totally asynchronous, event-based and does not use threads [4]. It instead uses the `fork()` system call for creating the master and workers processes. Hence problems arise because threads are not inherited across forks<sup>1</sup> The lazy initialization of the distributed driver in the `libsflr` workarounds this issue.

## 2 Communication protocol

HTTP requests for the `SF1` are required to specify in the URI the controller and the action parameters for the underlying `SF1` request, as in the sample Listing 1. The HTTP response body will then contain the `SF1` response in the specified format.

## 3 Configuration

The Listing 2 on the next page contains a sample configuration snippet containing the follow directives:

- `underscores_in_headers` [on|off] Nginx directive enabling HTTP headers in nonstandard format
- `rewrite` enables URI rewriting in order to replace `/sf1r/controller/action` to `/controller/action` as required by the driver<sup>2</sup>
- `sf1r` enables the `sf1r-module`
- `sf1r_addr host:port[,host:port]` [single|distributed] defines the target host; if the flag `distributed` is used, it is possible to define multiple hosts
- `sf1r_poolSize n` defines the connection pool initial size
- `sf1r_poolResize` [on|off] enables the pool auto-resize

<sup>1</sup>It is discouraged to use `fork()` and threads within the same program (see the glibc documentation at <http://www.imodulo.com/gnu/glibc/Threads-and-Fork.html>).

<sup>2</sup>See Nginx `HttpRewriteModule` [2].

Listing 2: Sample configuration

```
underscores_in_headers on;

location /sf1r/ {
    rewrite ^/sf1r(/.*)$ $1 break;

    sf1r;
    sf1r_addr server1:2181,server2:port2 distributed;
    sf1r_poolSize 5;
    sf1r_poolResize on;
    sf1r_poolMaxSize 10;
    sf1r_broadcast ^test/\w+$/;
    sf1r_broadcast ^recommend/visit_item$/;

    more_set_headers 'Access-Control-Allow-Origin: *';
    more_set_headers 'Access-Control-Allow-Methods: POST,
        GET, PUT, DELETE, OPTIONS';
    more_set_headers 'Access-Control-Allow-Headers: CONTENT-
        TYPE';
    more_set_headers 'Access-Control-Max-Age: 1728000';
    more_set_headers 'Access-Control-Allow-Credentials:
        false';
}
```

- `sf1r_poolMaxSize n` defines the maximum number of connections for the pool if the `sf1r_poolResize` has been set
- `sf1r_broadcast regex` define URI pattern for broadcasted requests
- `more_set_headers` additional HTTP response headers needed in order to support Ajax requests<sup>3</sup>

## 4 References

### References

- [1] libsf1r documentation, [git@izensoft.cn:izenelib.git](http://git@izensoft.cn:izenelib.git)
- [2] Nginx Wiki, <http://wiki.nginx.org/Main>
- [3] Evan Miller, *Emiller's Guide To Nginx Module Development*, <http://www.evanmiller.org/nginx-modules-guide.htm>
- [4] Joshua Zhou, *Nginx Internals Talk in Guangzhou, China*, <http://www.slideshare.net/joshzhu/nginx-internals>
- [5] Google Logging Library documentation, <http://google-glog.googlecode.com/svn/trunk/doc/glog.html>

---

<sup>3</sup>See third party `HttpHeadersMoreModule` [2].