

Nginx Notes

Paolo D'Apice

iZENESoft (Shanghai) Co. Ltd.

February 13, 2012



What's Nginx

A definition



Apache is like Microsoft Word, it has a million options but you only need six.

Nginx does those six things, and it does five of them 50 times faster than Apache.

Chris Lea



What's Nginx

Another definition



Batman is fast. Nginx is fast.

Batman fights crime. Nginx fights wasted CPU cycles and memory leaks.

Batman performs well under pressure. Nginx, for its part, excels under heavy server loads.

But Batman would be almost nothing without the Batman utility belt. Instead of a utility belt, Nginx has a module chain.

Evan Miller



How it works

Modules

Module types:

- handlers** process a request and produce an output

- filters** manipulate the output produced by a handler

- load-balancers** forward a request choosing between multiple backend servers



Components

1. configuration structs
2. directives
3. context
4. definition
5. installation



Components

Configuration structs

- ▶ up to three
 - ▶ main
 - ▶ server
 - ▶ location
- ▶ populated by module directives



Components

Directives

- ▶ array of `ngx_command_t` structs
- ▶ used for module configuration



Components

Context

- ▶ an `ngx_http_module_t` struct
- ▶ contains function references for creating and merging configurations



Components

Definition

- ▶ an `ngx_module_t` struct named `ngx_<module name>_module`
- ▶ contains references to context, directives and other callbacks



Components

Installation

- ▶ depends on the module type
- ▶ see later ...



Handlers

Typically do four things:

1. get location configuration
2. generate appropriate response
3. send response header
4. send response body



Handlers

Upstream Handlers (Proxy)

- ▶ do little "real work"
- ▶ callbacks invoked when the upstream is ready



Handlers

Installation

- ▶ define the directive that enables the module
- ▶ get the "core" struct for the location
- ▶ assign a handler to it



Filters

Manipulate responses generated by handlers

- ▶ header filters
- ▶ body filters



Filters

Header filters

Three basic steps:

1. decide whether to operate on the response
2. operate on the response
3. call the next filter



Filters

Body filters

- ▶ operate only on one buffer (chain link) at a time
 - ▶ overwrite
 - ▶ replace
 - ▶ insert new buffer before/after
- ▶ support incomplete buffer chain
- ▶ no nice high-level API



Filters

Installation

- ▶ post-configuration step
- ▶ insert the filter on the top of filter chain (LIFO)



Load-Balancer

Decide which backend server will receive a particular request

1. enabling configuration directive
2. registration function
3. upstream initialization function
4. peer initialization function
5. load balancing function
6. peer release function



Load-Balancer

enabling directive

- ▶ should have the `NGX_UPS_CONF` flag set
- ▶ provide a pointer to a *registration function*



Load-Balancer

registration function

- ▶ define options within the `upstream` configuration
- ▶ register an *upstream initialization function*



Load-Balancer

upstream initialization function

- ▶ resolves the server names
- ▶ allocates sockets
- ▶ sets a callback to the *peer initialization function*



Load-Balancer

peer initialization function

- ▶ called *once* per request
- ▶ sets up data structures used by the *load balancing function*
 - ▶ struct named `ngx_http_upstream_<module>_peer_data_t`
- ▶ sets up two callbacks
 - `get` the *load-balancing function*
 - `free` the *peer release function*



Load-Balancer

load balancing function

- ▶ called *at least* per request
- ▶ defines the load-balancing policy using sockets



Load-Balancer

peer release function

- ▶ operates after the upstream connection take place
- ▶ at least must set `pc->tries` to zero



Writing Nginx Modules

- ▶ no or little documentation
 - ▶ mostly in Russian or Chinese
 - ▶ *self documented* source code
- ▶ learn by example



Writing Nginx Modules

- ▶ write you code
- ▶ write a `config` file
- ▶ compile Nginx with your module



Writing Nginx Modules

- ▶ take a look at the hello-module
 - ▶ autotools adapted to CMake
 - ▶ C++ support¹
 - ▶ Test::Nginx Perl framework
- ▶ stick to the naming standards

¹Nginx is written in plain C



References



Nginx Wiki

<http://wiki.nginx.org/Main>



Evan Miller's *definitive* guide

<http://www.evanmiller.org/nginx-modules-guide.htm>



Antoine Bonavita's blog about Nginx

<http://www.nginx-discovery.com>



More References



Taobao's Tengine

<http://tengine.taobao.org>



agentz's blog

<http://agentzh.org>



chaoslawful's blog

<http://chaoslawful.iteye.com>



Joshua Zhu's blog

<http://blog.zhuzhaoyuan.com>

